

# Review

May 2023

---

## Slides 1:

- Levels of parallelism
- Amdahl's law

---

## Slides 2:

- Pipelining
  - Dependencies
  - Loop fusion
  - Branch prediction
  - Optimizations

---

## Slides 3:

- Caches
  - Cache hit
  - Cache misses: Capacity, Conflict, Compulsory
  - Temporal locality
  - Spatial locality
  - Cache organization: Direct-mapped, Fully associative, Set-associative
  - Cache organization: Multi-level caches - why?

---

## Slides 4:

- Improving cache performance
  - Cache blocking
  - Examples: Matrix traversal, Jacobi

---

**Slides 5:**

- Improving cache performance
  - Examples: Matrix transpose, Matric-Vector, Matrix-Matrix multiplication

---

**Slides 6:**

- Parallelism: Vector instructions
  - Trust the compiler
  - Assembly language
  - OpenMP SIMD
  - Vector instrincs
- Programmer intervention: `__attribute__` `__restrict__`

---

**Slides 7 and 8:**

To know:

- Processes
- Threads
- Hardware and Software threads
- SMT (Simultaneous Multi-Threading) / Hyperthreading - what does it mean?
- Address space of a process
- Role of the Operating System

Multicore programming:

- Cores - do they share caches?
- False sharing - does it affect performance?
- Coherence - what is it?
- Synchronization

---

**Slides 9:**

- Coherency and Consistency - what is the difference?

---

**Slides 10 and 11:**

- OpenMP
- Scheduling - static, dynamic, guided, auto
- Nested parallelism

Examples:

- Matrix-matrix multiplication
- Jacobi
- SparseMV
- Gauss-Seidel

---

**Slides 12: OpenMP: Tasks**

- Example: Graph traversal

---

**Slides 13, 14, 15, 16 and 17: To know:**

- What does it mean to say - this MPI program is launched on 4 processes?
- Can the 4 processes be located on the same machine?
  - If yes, do they share memory i.e. can a variable declared in one process be accessed by another without using any MPI calls?
  - If no, how do they communicate data?
  - Do they share physical resources like memory and network?

Important functions: MPI\_Init MPI\_Comm\_size MPI\_Comm\_rank MPI\_Send MPI\_Recv MPI\_Finalize

MPI\_Send and MPI\_Recv are blocking calls. Read about non-blocking calls. Example: MPI\_Isend() MPI\_Irecv() Why are non-blocking calls useful?

Other important functions: MPI\_Barrier MPI\_Bcast MPI\_Reduce MPI\_Allreduce MPI\_Scatter MPI\_Gather MPI\_Allgather

Questions to ask:

- Can one implement MPI\_Barrier or MPI\_Bcast using MPI\_Send and MPI\_Recv?
- How about other functions?

Examples:

- Jacobi
  - Data distribution
  - How did we avoid deadlock?
  - How did we overlap computation and communication?
- Graph traversal
- Jacobi 2D distribution
  - MPI\_Type\_Vector
  - MPI\_Cart\_Create
- Matrix product: 1D distribution

---

**Slides 18, 19 and 20:**

- We discuss in detail the implementation of Molecular Dynamics using MPI.
- We explore several design choices.

---

**Slides 21:**

- MPI + OpenMP
- One-sided communication. Examples:
  - Jacobi 1D tiling
  - Distributed table
- We have slides on other programming paradigms (is not part of exam syllabus, but is useful to know):
  - PGAS
  - Charm++

---

**GPU:**

Core ideas:

- What happened to some of the components?
- Fuse fetch-decode logic - what is the consequence of that?
- Store many contexts
- Threads, Blocks, Grid : CUDA cores, warps, SMs
- How is memory organized?
- What is memory coalescing?

Thrust, Streams and Dynamic parallelism - not part of exam syllabus, but useful to know.

---

### **GPU OpenMP:**

using OpenMP

- omp target
- map
- Jacobi example

Principal Component Analysis example - not part of exam syllabus, but useful to understand how to use OpenMP in GPU better.

---

**MPI-IO:** Each process makes independent request; Collective I/O: All processes make a single request.

---

**openmp\_2 in General Resources:** OpenMP example: Merge Sort - not part of exam syllabus, but useful to know.

---