

CS420 – Lectures 12

Raghavendra Kanakagiri
Slides: Marc Snir

Spring 2023



Barriers

```
double a[N], b[N], c[N];
#pragma omp parallel
{
    int tid = omp_get_thread_num();
    if (tid == 0) {
        for (int i = 0; i < N; i++) a[i] = 1.0;
    }
    #pragma omp barrier
    #pragma omp for
    for (int i = 0; i < N; i++) {
        b[i] = func(a, i);
    }
    // implicit barrier
    #pragma omp for nowait
    for (int i = 0; i < N; i++) {
        c[i] = func(a, i);
    }
    // no barrier
    a[tid] = 0.0;
}
// implicit barrier
```

Data Sharing

```
double a[N];
int main()
{
    int x = 3;
    #pragma omp parallel
    {
        double y;
        // some work
    }
    return 0;
}
// x is shared, y is private
// private -> each thread gets its own copy
```

Data Sharing

```
double a[N];
int main()
{
    int sum = 0;
    #pragma omp parallel for private(sum)
    for (int i = 0; i < N; i++)
    {
        sum += a[i];
    }
    printf("sum = %d\n", sum);
    return 0;
}
// what is wrong here?
```

Data Sharing

```
double a[N];
int main()
{
    int sum = 0;
    #pragma omp parallel for private(sum)
    for (int i = 0; i < N; i++)
    {
        sum += a[i];
    }
    printf("sum = %d\n", sum);
    return 0;
}
// sum in each thread is uninitialized
// print sum = 0
```

Data Sharing

```
double a[N];
int main()
{
    int sum = 0;
    #pragma omp parallel private(sum)
    {
        #pragma omp for
        for (int i = 0; i < N; i++)
        {
            sum += a[i];
        }
        printf("sum = %d\n", sum);
    }
    printf("sum = %d\n", sum);
    return 0;
}
```

- Parallel loops are convenient for nice iteration domains, but not for irregular computations where it is not clear upfront what tasks need to be generated.
- The *task* construct helps for this purpose.

Within a parallel section

```
#pragma omp task  
{...}
```

will start a task that can execute on any of the available threads; the calling task may continue executing in parallel with the newly created task.

Tasks

- `task`: spawns a task that can execute separately
- `taskwait`: wait for all spawned tasks to complete before continuing
- `shared`: parent's variable shared with child

Tasks

```
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp task
        func1();
        #pragma omp task
        func2();
        #pragma omp taskwait
        #pragma omp task
        func3();
    }
}
```

// func3() can execute only after func1() and func2() have completed

Tasks

```
#pragma omp parallel shared(x) private(y)
{
    #pragma omp task
    {
        int z;
        func(x, y, z);
    }
}
// s is shared, y is firstprivate, z is private
```

A terrible example: Fibonacci

$\text{fib}(0) = 0; \text{fib}(1)=1;$
 $\text{fib}(n)=\text{fib}(n-1)+\text{fib}(n-2)$

```
int fib(int n) {  
    int i, j;  
    if (n<2) return n;  
    else {  
        #pragma omp task shared(i)  
        i=fib(n-1);  
        #pragma omp task shared(j)  
        j=fib(n-2);  
        #pragma omp taskwait  
        return i+j;  
    }  
}
```

- task: spawns a task that can execute separately
- taskwait: wait for all spawned tasks to complete before continuing
- shared: parent's variable shared with child


Why terrible?

- Can be computed in constant time:

$$\text{fib}(n) = \frac{1}{\sqrt{5}} \left(\left(\frac{1+\sqrt{5}}{2} \right)^n - \left(\frac{1-\sqrt{5}}{2} \right)^n \right) \approx \frac{1}{\sqrt{5}} \left(\left(\frac{1+\sqrt{5}}{2} \right)^n \right)$$


- Can be computed in linear time using the linear recursion
- Number of tasks spawned by the parallel algorithm is $\text{ntasks}(n) = \text{ntasks}(n-1) + \text{ntasks}(n-2)$; i.e., $\text{ntasks}(n) = \text{fib}(n)$. Exponential amount of compute work!

Example: Graph traversal




Mark all the nodes that can be reached from node 1


Example: Graph traversal



Mark all the nodes that can be reached from node 1



Adjacency list representation



1	2	3	
2	1	3	6
3	1	2	4
4	3	5	
5	4	6	
6	2	5	7
7	6		
8	9		
9	8		

Parallel traversal

```
\\ structure for node
typedef struct {
    int visited; \\ mark for visited node
    int numneighbors; \\ number of neighbors (degree)
    int neighbors[]; \\ array of neighbor ids
} Node;
```

```
Node * graph[N]; \\ array of pointers to nodes
```

```
void visit(int i) {
    int j,k,mark;
    for(j=0; j<graph[i]->numneighbors; j++) {
        k = graph[i]->neighbors[j];
        #pragma omp atomic
        mark = graph[k]->visited++;
        if(mark==0)
            #pragma omp task
            visit(k);
    }
}
}
```



```
int main() {  
    #pragma omp parallel \\ need to start all threads  
    #pragma omp single \\ need to call only once  
        visit(0);  
}
```

Task Dependences

True dependence (aka RAW, aka *flow dependence*)

```
int main() {  
    int x = 1;  
    #pragma omp parallel  
        #pragma omp single {  
            #pragma omp task shared(x) depend(out: x)  
                x = 2;  
            #pragma omp task shared(x) depend(in: x)  
                printf("x = %d\n", x); } return 0;  
}
```

Will print x=2

Task Dependences

Anti-dependence (aka WAR)

```
int main() {  
    int x = 1;  
    #pragma omp parallel  
        #pragma omp single  
        {  
            #pragma omp task shared(x) depend(in: x)  
            printf("x = %d\n", x);  
            #pragma omp task shared(x) depend(out: x)  
            x = 2;  
        }  
    return 0;  
}
```

Will print x=1

Task Dependences

Output-dependence (aka WAW)

```
int main() {  
    int x;  
    #pragma omp parallel  
    #pragma omp single  
    {  
        #pragma omp task shared(x) depend(out: x)  
        x = 1;  
        #pragma omp task shared(x) depend(out: x)  
        x = 2;  
        #pragma omp taskwait  
        printf("x = %d\n", x);  
    }  
    return 0;  
}
```

Will print x=2

If a dependence exists then tasks are executed in the order they were spawned.

Not a dependency

RAR

```
int main() {
    int x = 1;
    #pragma omp parallel
    #pragma omp single
    {
        #pragma omp task shared(x) depend(in: x)
        printf("x = %d\n", x);
        #pragma omp task shared(x) depend(in: x)
        x = 2;
    }
    return 0;
}
```

Can print x=1 or x=2